

Artificial Intelligence

1 What is an AI?

There are various definitions of an AI, ranging from thinking humanly and rationally to acting humanly and rationally. The *turing test*, is test in which a human interrogator interacts with a machine, sending it messages back and forth, and a machine passes if it fools the human into thinking that the messages are being sent to them by a human. For this a computer needs: **natural language processing, knowledge representation, automated reasoning and machine learning**. To pass the *total turing test* a computer would additionally need **computer vision and robotics**.

1.1 Intelligent Agents

An **agent** is just something that acts and a **rational agent** is one that acts so as to achieve the best outcome or, when there is uncertainty, the best expected outcome. **Percept** means the agent's perceptual inputs at any given time, and a **percept sequence** is the complete history of everything the agent has ever perceived. The **agent function** is an abstract mathematical description that maps a given percept to an action; an **agent program** is a concrete implementation of the agent function, running within some physical system. *It is better to design a performance measure according to what one wants in an environment, then how one wants an agent to behave.*

The proper definition of a rational agent is *for each possible percept sequence, a rational agent should select an action that is expected to maximize its performance measure, given the evidence provided by the percept sequence and whatever built-in knowledge the agent has*. An **omniscient** agent knows the actual outcomes of its actions.

1.2 The Nature of Environments

An environment is defined as **PEAS**: performance measure, environment, actuators and sensors. There are different types of environments, namely:

- **Observable vs Partially-Observable**: it is observable when the agents' sensors have complete access to the environment's state at all times

- **Single agent vs Multi agent:** there could be multiple agents in an environment. There is also a question of what must be considered an agent. This gives way to the concept of **competitive** vs **cooperative** environments.
- **Deterministic vs Stochastic:** If the next state can be completely determined by the current state and the action executed by the agent, then it is deterministic; and stochastic otherwise. An environment is **uncertain** if it is not fully observable or not deterministic. Note that a **non-deterministic** environment is one where each action is characterized by its possible outcomes, but no probabilities are attached to them.
- **Episodic vs Sequential:** In an episodic environment the agent's experience is divided into atomic episodes. The next episode doesn't depend on the action taken in the previous episode.
- **Static vs Dynamic:** If an environment can change when an agent is deliberating, then it's dynamic, and is static otherwise. If the environment doesn't change when deliberating but the performance score does, then we call it **semi-dynamic**.
- **Discrete vs Continuous:** The distinction here applies to the state of the environment, the way time is handled and the percepts and actions of the agent. For example, chess having a discrete set of states; the same doesn't apply for taxi driving.
- **Known vs Unknown:** This applies to the agent's state of knowledge about the "laws of physics" of the environment. Note that it's possible that a known environment is partially observable like solitaire. Conversely, an environment can also be unknown and fully observable, like in a video game, one can see the state but one doesn't know the control until one tries to play.

1.3 The Structure of Agents

There are four basic types of agent programs:

- **Simple Reflex Agents:** Agents that select the current action based on the current precepts and ignoring the rest of the precept history. It is also important to note that these types of agents are usually implemented in a fully-observable environment.
- **Model-Based Reflex Agents:** The best way to handle a partially observable environment is to keep some sort of an internal representation of the aspects of the environment not currently observable. Therefore, an agent should have some sort of knowledge about how the world works and the agents who have said knowledge are called model-based agents.

- **Goal-Based Agents:** These types of agents consider how close they get to a goal, in addition to have a model of how their environment works. These types of agents are also quite flexible as they can update their actions on-the-fly depending on their goals and the feedback they get from the environment.
- **Utility-Based Agents:** Since the previous model does not differentiate between how it gets to its goal, and which state would make it more happy, it is not efficient. Therefore a utility function is needed to determine just that i.e. it is an internalization of its performance measure. The previous model will also fail when there are conflicting goals or when there are several goals the agent can aim for, none of which can be achieved with certainty; in both cases, a utility function can dictate which action to take to maximize expected utility.

There are different ways an agent can represent the world around it:

- **Atomic:** Each state of the world is indivisible: it has no internal structure.
- **Factored:** Each state is split up into a fixed set of variables and attributes, each of which can have a value. Uncertainty can also be represented in this representation.
- **Structured:** This type of representation has objects and their relationship with each other.

2 Problem Solving by Searching

Goal formulation, based on the current situation and the agent's performance measure, is the first step in problem solving. **Problem formulation** is the process of deciding what actions and states to consider, given a goal. In general, *an agent with several immediate options of unknown value can decide what to do by first examining future actions that eventually lead to states of known value.* The process of looking for a sequence of actions that reaches the goal is called **search**. A search algorithm takes a problem as input and returns a **solution** in the form of an action sequence. Once a solution is found, the actions it recommends can be carried out. This is called the **execution** phase. Note while the agent is executing, it *ignores its percepts* when choosing its actions because it knows in advance what they will be.

Together, the **initial state**, **actions** and **transition model** implicitly define the **state space** of the problem - the set of all states reachable from the initial state by any sequence of actions. A **solution** to a problem is an action sequence that leads from the initial state to a goal state. Solution quality is measured by the path cost function, and an **optimal solution** has the lowest

path cost among all solutions. The process of removing detail from a representation is called **abstraction**.

It is quite important to distinguish between a node and a state: a node is a bookkeeping data structure used to represent the search tree. A state corresponds to a configuration of the world. Furthermore, two different states can contain the same world state if that state is generated via two different search paths. We can evaluate the performance of an algorithm in four ways:

- **Completeness:** Is the algorithms guaranteed to find a solution if there is one?
- **Optimality:** Does the strategy find the optimal solution?
- **Time complexity:** How long does it take to find a solution?
- **Space complexity:** How much memory is needed to perform the search?

In AI, the graph is often represented implicitly by the initial state, actions and transition model and is frequently infinite. Complexity is expressed in terms of three quantities: b , the **branching factor** or maximum number of successors of any node; d , the **depth** of the shallowest goal node; and m , the maximum length of any path in the state space.

2.1 Uninformed Search Strategies

2.1.1 Breadth-First Search

This is a simple strategy in which all the nodes are expanded at a given depth in the search tree before any nodes at the next level are expanded. The algorithm is given in figure 1.

BFS is *complete* - if the shallowest goal node is at some finite depth BFS will eventually find it after generating all the shallower nodes. Note that the *shallowest* is not always the *optimal* one. BFS is only optimal if the path cost is a non-decreasing function of the depth of the node. The time complexity of BFS is:

$$b + b^2 + b^3 + \dots + b^d = O(b^d) \quad (1)$$

For the space complexity there will be $O(b^{d-1})$ nodes in the explored set and $O(b^d)$ nodes in the frontier, giving space complexity as $O(b^d)$. *Memory requirements are a bigger problem for BFS than execution time and time is still a major factor. Exponential-complexity search problems cannot be solved by uninformed methods for any but the smallest instances.*

2.1.2 Uniform-Cost Search

This is similar to BFS, but it expands the node with the lowest path cost, the goal test is applied to a node *not when it's generated, but when it's chosen for expansion* and a test is added in case a better path is found to a node currently

```

function BREADTH-FIRST-SEARCH(problem) returns a solution, or failure
  node  $\leftarrow$  a node with STATE = problem.INITIAL-STATE, PATH-COST = 0
  if problem.GOAL-TEST(node.STATE) then return SOLUTION(node)
  frontier  $\leftarrow$  a FIFO queue with node as the only element
  explored  $\leftarrow$  an empty set
  loop do
    if EMPTY?(frontier) then return failure
    node  $\leftarrow$  POP(frontier) /* chooses the shallowest node in frontier */
    add node.STATE to explored
    for each action in problem.ACTIONS(node.STATE) do
      child  $\leftarrow$  CHILD-NODE(problem, node, action)
      if child.STATE is not in explored or frontier then
        if problem.GOAL-TEST(child.STATE) then return SOLUTION(child)
        frontier  $\leftarrow$  INSERT(child, frontier)

```

Figure 1: Breadth-First Search.

on the frontier. The algorithm is given in figure 2. First, we note that whenever this algorithm selects a node for expansion, the optimal path to that node is found and second, paths never get shorter as nodes are added. Thus, *uniform-cost search expands nodes in order of their optimal path cost*. This search is complete given that the cost of every step exceeds some small positive constant ϵ . The space and time complexity is $O(b^{1+\lfloor C^*/\epsilon \rfloor})$, where C^* is the cost of the optimal solution.

2.1.3 Depth-First Search

Depth-First Search always expands upon the deepest node in the frontier, this is accomplished using a LIFO queue. The graph-search version of the algorithm is complete in finite spaces but the tree-search version could potentially fall into an infinite loop. Although, it must be noted that both versions fail if there is an infinite state space, with an infinite non-goal path e.g. Knuth's 4 problem. Both versions are also not optimal. The time complexity of DFS is $O(b^m)$ and space complexity is $O(bm)$.

2.1.4 Depth-Limited Search

The problem of DFS failing in infinite spaces can be solved if we apply a limit ℓ to the depth we expand up till. Although, it introduces incompleteness if we choose $\ell < d$ and non-optimality if $\ell > d$. Its time complexity is $O(b^\ell)$ and its space complexity is $O(b\ell)$. Its pseudocode is shown in figure 4.

2.1.5 Iterative Deepening Search

The algorithm of IDS is shown in figure 5. Its space complexity is $O(bd)$. Like BFS, IDS is complete if the branching factor is finite and it is optimal when the

```

function UNIFORM-COST-SEARCH(problem) returns a solution, or failure
  node ← a node with STATE = problem.INITIAL-STATE, PATH-COST = 0
  frontier ← a priority queue ordered by PATH-COST, with node as the only element
  explored ← an empty set
  loop do
    if EMPTY?(frontier) then return failure
    node ← POP(frontier) /* chooses the lowest-cost node in frontier */
    if problem.GOAL-TEST(node.STATE) then return SOLUTION(node)
    add node.STATE to explored
    for each action in problem.ACTIONS(node.STATE) do
      child ← CHILD-NODE(problem, node, action)
      if child.STATE is not in explored or frontier then
        frontier ← INSERT(child, frontier)
      else if child.STATE is in frontier with higher PATH-COST then
        replace that frontier node with child

```

Figure 2: Uniform-Cost Search.

A recursive implementation of DFS:^[5]

```

1 procedure DFS(G, v):
2   label v as discovered
3   for all edges from v to w in G.adjacentEdges(v) do
4     if vertex w is not labeled as discovered then
5       recursively call DFS(G, w)

```

A non-recursive implementation of DFS:^[6]

```

1 procedure DFS-iterative(G, v):
2   let S be a stack
3   S.push(v)
4   while S is not empty
5     v = S.pop()
6     if v is not labeled as discovered:
7       label v as discovered
8       for all edges from v to w in G.adjacentEdges(v) do
9         S.push(w)

```

Figure 3: Depth-First Search.

```

function DEPTH-LIMITED-SEARCH(problem, limit) returns a solution, or failure/cutoff
return RECURSIVE-DLS(MAKE-NODE(problem.INITIAL-STATE), problem, limit)

function RECURSIVE-DLS(node, problem, limit) returns a solution, or failure/cutoff
if problem.GOAL-TEST(node.STATE) then return SOLUTION(node)
else if limit = 0 then return cutoff
else
    cutoff_occurred? ← false
    for each action in problem.ACTIONS(node.STATE) do
        child ← CHILD-NODE(problem, node, action)
        result ← RECURSIVE-DLS(child, problem, limit − 1)
        if result = cutoff then cutoff_occurred? ← true
        else if result ≠ failure then return result
    if cutoff_occurred? then return cutoff else return failure

```

Figure 4: Depth-Limited Search.

```

function ITERATIVE-DEEPENING-SEARCH(problem) returns a solution, or failure
for depth = 0 to ∞ do
    result ← DEPTH-LIMITED-SEARCH(problem, depth)
    if result ≠ cutoff then return result

```

Figure 5: Iterative Deepening Search.

path cost is a non-decreasing function of the depth. The time complexity is:

$$(d)b + (d-1)b^2 + \dots + (d-d+1)b^d = O(b^d) \quad (2)$$

In general, IDS is preferred if the search space is large and the depth of the solution is not known.

2.1.6 Bidirectional Search

Bidirectional search applies the idea of starting two searches: one from the initial state and one from the goal state, hoping that they meet in the middle; with the motivation that $b^{d/2} + b^{d/2} < b^d$. Thus the space and time complexity for this algorithm is $O(b^{d/2})$.

2.1.7 Comparison of Uninformed Search Strategies

Figure 6 shows comparisons for tree search versions of the algorithms discussed. For graph searches, the main difference is that DFS is complete for finite state spaces and that the space and the time complexities are bounded by size of the state space.

Criterion	Breadth-First	Uniform-Cost	Depth-First	Depth-Limited	Iterative Deepening	Bidirectional (if applicable)
Complete?	Yes ^a	Yes ^{a,b}	No	No	Yes ^a	Yes ^{a,d}
Time	$O(b^d)$	$O(b^{1+\lceil C^*/\epsilon \rceil})$	$O(b^m)$	$O(b^\ell)$	$O(b^d)$	$O(b^{d/2})$
Space	$O(b^d)$	$O(b^{1+\lceil C^*/\epsilon \rceil})$	$O(bm)$	$O(b\ell)$	$O(bd)$	$O(b^{d/2})$
Optimal?	Yes ^c	Yes	No	No	Yes ^c	Yes ^{c,d}

Figure 6: Comparison of uninformed search strategies. b is the branching factor; d is the depth of the shallowest solution; m is the maximum depth of the search tree; ℓ is the depth limit. Superscript caveats are as follows: ^a complete if b is finite; ^b complete if step costs $\geq \epsilon$ for positive ϵ ; ^c optimal if step costs are all identical; ^d if both directions use breadth-first search.

3 Informed Search Algorithms

An informed search strategy is one that uses problem-specific knowledge beyond the definition of the problem itself. The general approach considered is an instance of a general graph or tree search in which a node is chosen for expansion based on its **evaluation function**, $f(n)$. Most best-first algorithms include as a component of f a **heuristic function**, denoted by $h(n)$, which is an estimated cost of the cheapest path from the state at node n to a goal state. We use the constraint that if n is the goal node then $h(n) = 0$.

3.1 Greedy Best-First Search

This algorithm uses $f(n) = h(n)$. Due to its greedy nature it is not optimal and is also incomplete even in a finite state space, like DFS. The graph version of this algorithm is complete in finite spaces, but not infinite ones. The worse case space and time complexity for the tree version is $O(b^m)$, where m is the maximum depth of the search tree.

3.2 A* Search

This search evaluates nodes by using:

$$f(n) = g(n) + h(n) \quad (3)$$

Since $g(n)$ gives the path cost from the start node to node n , and $h(n)$ is the estimated cost of the cheapest path from n to the goal, $f(n)$ is the estimated cost of the cheapest solution through n . Provided $h(n)$ satisfies certain conditions, it's both complete and optimal.

3.2.1 Conditions for Optimality: Admissibility and Consistency

The first condition we require is $h(n)$ be an **admissible heuristic**, meaning it is one that *never overestimates* the cost to reach the goal. Thus we have

that $f(n)$ never overestimates the true cost of a solution along the current path through n . The second condition required is **consistency**, or **monotonicity** for applications of A^* to graph search. A heuristic $h(n)$ is consistent if:

$$h(n) \leq c(n, a, n') + h(n') \quad (4)$$

Where n' is the successor of n generated by action a . This is a form of the general **triangle inequality**, which says that each side of a triangle can't be longer than the sum of the other two sides. Here the triangle is formed by n, n' and the goal G_n closest to n .

3.2.2 Optimality of A^*

We know that *the tree-search version of A^* is optimal if $h(n)$ is admissible, while the graph-search version is optimal if $h(n)$ is consistent*. First let us establish: *if $h(n)$ is consistent then the values of $f(n)$ along any path are non-decreasing*. Suppose n' is a successor of n ; then $g(n') = g(n) + c(n, a, n')$ thus we have:

$$f(n') = g(n') + h(n') = g(n) + c(n, a, n') + h(n') \geq g(n) + h(n) = f(n) \quad (5)$$

The next step would be to prove that *whenever A^* selects a node for expansion, the optimal path to that node has been found*. The fact that f costs are non-decreasing along any path allows us to draw **contours** in the state space. If C^* is the cost of the optimal solution path, then:

- A^* first expands all nodes with $f(n) < C^*$.
- A^* might then expand some of the nodes right on the “goal contour” (where $f(n) = C^*$) before selecting the goal node.

Completeness requires that there be only finitely many nodes with cost less than or equal to C^* , a condition only met if all step costs exceed some finite ϵ and if b is finite. One final observation is A^* is **optimally efficient** for any given consistent heuristic i.e. no other optimal algorithm is guaranteed to expand fewer nodes than A^* because any algorithm that doesn't expand all nodes with $f(n) < C^*$ runs the risk of missing the optimal solution. The space complexity is $O(b^d)$ and the time complexity is $O(b^\Delta)$, with $\Delta \equiv h^* - h$, where h^* is the actual cost of getting from the root to the goal; and this is the formula for **absolute error**. For constant step costs we have $O(b^{\epsilon d})$, with $\epsilon \equiv (h^* - h)/h^*$ which is the **relative error**.

3.3 Heuristic Functions

There are a few good heuristics used for the 8-puzzle like the number of misplaced tiles and the **manhattan distance** which is the sum of horizontal and vertical distances. One way to characterize the quality of a heuristic is the **effective branching factor**, b^* , which is conventionally defined as an average number of nodes revisited of the current iteration N as compared to the previous

iteration $N - 1$. If the total number of nodes generated by A* for a particular problem is N , and the solution depth is d , then b^* is the branching factor that a uniform tree of depth d would have in order to contain $N + 1$ nodes. Thus:

$$N + 1 = 1 + b^* + (b^*)^2 + \dots + (b^*)^d \quad (6)$$

If \forall nodes $n, h_2(n) \geq h_1(n) \implies h_2$ **dominates** h_1 .

3.3.1 Relaxed Problems

A problem with fewer restrictions on the actions is called a **relaxed problem**. The state-space graph of the relaxed problem is a *supergraph* of the original state space because the removal of restrictions creates added edges in the graph. Because edges are added to the state-space, any optimal solution in the original problem is also a solution in the relaxed problem, but the relaxed problems may have better solutions. Hence, *the cost of an optimal solution to a relaxed problem is an admissible heuristic for the original problem*. Furthermore, because the derived heuristic is an exact cost for the relaxed problem, it must obey the triangle inequality and is therefore consistent. If the relaxed problem is hard to solve, then the values of the corresponding heuristic will be expensive to compute. Since a collection of admissible heuristics are obtained, we can use an appropriate heuristic on each node individually using:

$$h(n) = \max\{h_1(n), \dots, h_m(n)\} \quad (7)$$

3.4 Local Search Algorithms

These algorithms operate using a single **current node** and generally move only to neighbors of that node and the paths followed aren't retained. Thus they have the advantages:

- They use very little memory - usually a constant amount
- They can often find reasonable solutions in large or infinite state spaces

3.4.1 Hill-Climbing Search

The algorithm is shown in figure 7. It is comparative to “trying to find the top of Mount Everest in a thick fog while suffering from amnesia”. It is also called **greedy local search**. Hill-Climbing gets stuck in **local maxima/minima**, **ridges** and **plateaux**.

4 Adversarial Search

A **utility function** is defined as the final numeric value for a game that ends in terminal state s for a player p . A **zero-sum game** is defined as one where the total payoff to every player is the same in every instance of the game.

```

function HILL-CLIMBING(problem) returns a solution state
  inputs: problem, a problem
  local variables: current, a node
                   next, a node

  current ← MAKE-NODE(INITIAL-STATE[problem])
  loop do
    next ← a highest-valued successor of current
    if VALUE[next] < VALUE[current] then return current
    current ← next
  end

```

Figure 7: Hill-Climbing Search.

```

function MINIMAX-DECISION(state) returns an action
  return  $\arg \max_{a \in \text{ACTIONS}(s)} \text{MIN-VALUE}(\text{RESULT}(\text{state}, a))$ 

```

```

function MAX-VALUE(state) returns a utility value
  if TERMINAL-TEST(state) then return UTILITY(state)
  v ←  $-\infty$ 
  for each a in ACTIONS(state) do
    v ← MAX(v, MIN-VALUE(RESULT(s, a)))
  return v

```

```

function MIN-VALUE(state) returns a utility value
  if TERMINAL-TEST(state) then return UTILITY(state)
  v ←  $\infty$ 
  for each a in ACTIONS(state) do
    v ← MIN(v, MAX-VALUE(RESULT(s, a)))
  return v

```

Figure 8: The Minimax Algorithm.

4.1 The Minimax Algorithm

The algorithm is shown in figure 8. The utility values of the leaf nodes are computed and the corresponding minimax values are backed up the tree recursively. The time complexity of this algorithm is $O(b^m)$, where m is the maximum depth of the tree and there are b legal moves at each point. The space complexity is $O(bm)$ for an implementation that generates all the actions at once and $O(m)$ for an implementation that generates actions one at a time.

4.2 Alpha-Beta Pruning

A way to improve the abysmal complexity of Minimax is to use alpha-beta pruning, which prunes away branches that can't possibly influence the final decision. α is the best choice we have found so far for Max and β is the best

```

function ALPHA-BETA-SEARCH(state) returns an action
   $v \leftarrow \text{MAX-VALUE}(\text{state}, -\infty, +\infty)$ 
  return the action in ACTIONS(state) with value v

```

```

function MAX-VALUE(state,  $\alpha$ ,  $\beta$ ) returns a utility value
  if TERMINAL-TEST(state) then return UTILITY(state)
   $v \leftarrow -\infty$ 
  for each a in ACTIONS(state) do
     $v \leftarrow \text{MAX}(v, \text{MIN-VALUE}(\text{RESULT}(s, a), \alpha, \beta))$ 
    if  $v \geq \beta$  then return v
     $\alpha \leftarrow \text{MAX}(\alpha, v)$ 
  return v

```

```

function MIN-VALUE(state,  $\alpha$ ,  $\beta$ ) returns a utility value
  if TERMINAL-TEST(state) then return UTILITY(state)
   $v \leftarrow +\infty$ 
  for each a in ACTIONS(state) do
     $v \leftarrow \text{MIN}(v, \text{MAX-VALUE}(\text{RESULT}(s, a), \alpha, \beta))$ 
    if  $v \leq \alpha$  then return v
     $\beta \leftarrow \text{MIN}(\beta, v)$ 
  return v

```

Figure 9: Alpha-Beta Pruning.

choice for Min. The algorithm is applied in figure 9. **Move ordering** can be applied to improve complexity to $O(b^{m/2})$ and the branching factor becomes \sqrt{b} . If the successors are examined at random and not best-first the complexity increases to $O(b^{3m/4})$. The best moves in a game are called **killer moves** and a killer move heuristic is used to find them. The hash table of previously seen moves is called a **transposition table**.

4.3 Evaluation Functions

An evaluation function returns an estimate of the expected utility from a given position. This is useful as sometimes we want to stop our search at a depth limit and evaluate the leaf nodes using this function. Most evaluation functions compute separate numerical contributions from each feature then combine them to find the total value. These types of functions are also called **weighted linear functions** because:

$$E(s) = w_1 f_1(s) + w_2 f_2(s) + \dots + w_n f_n(s) = \sum_{i=1}^n w_i f_i(s) \quad (8)$$

In **cutting off search** isTerminal and utility in minimax are replaced by cutoff and eval. In stochastic games, chance nodes are added in addition to min/max nodes. The minimax values are also replaced by **expected values**: the average over all possible outcomes of the chance nodes. This leads us to make **expecti-minimax** for games with chance nodes.

5 Machine Learning in Game Search

We can automate the fine design choices in an evaluation function by using machine learning. We can use **book learning** to learn a sequence of moves for important positions, like remembering what opening moves lead to what outcomes and remembering what moves led to a loss and avoiding them in the future. **Search control learning** is used to learn how to make search more efficient like order of move generation for $\alpha - \beta$ and learning a classifier to predict what depth we should search depending on the current state. Weights in evaluation functions can also be learned to make them agree with the true final utility.

5.1 Gradient Descent Learning

This is a type of *supervised learning* which is like the Hill-Climbing algorithm, in that it starts at some point \mathbf{w} in the weight space and moves to a neighboring lower point until we converge to the minimum possible loss. The hope that weights can be learned that closely approximate the true output. Each weight is updated according to:

$$w_i \leftarrow w_i - \alpha(z - t)f_i(s) \quad (9)$$

α is the **learning rate** which can be a constant or can decay over time. The problems with this type of learning are:

- Delayed reinforcement: reward maybe received after a few steps of applying the move
- Credit assignment: need to know which actions were responsible for the outcome

5.2 Temporal Difference Learning

Since supervised learning is for single-step prediction, TD is for multi-step prediction. The correctness of the prediction is not known until a few steps later, but intermediate steps can provide information about correctness. This is a type of supervised learning. **TDLeaf**(λ) combines TD and minimax to update evaluation function to reduce differences in rewards predicted at different levels of the search. The weight update rule applied is:

$$w_j \leftarrow w_j + \alpha \sum_{i=1}^{N-1} \frac{\partial r(s_i^l, w)}{\partial w_j} \sum_{m=1}^{N-1} \lambda^{m-i} d_m \quad (10)$$

Where $r(s_i^l, w)$ is the reward for of a state (which is found at max cut-off depth using minimax starting at s_i) using the weight w . And d_i is defined as difference between successive states:

$$d_i = r(s_{i+1}^l, w) - r(s_i^l, w) \quad (11)$$